



Inspire...Educate...Transform.

Engineering Big Data

Processing Frameworks on Clusters:
Map-Reduce

Dr. Sreerama KV Murthy
CEO, Quadratyx

August 8, 2015

Wake-Up Quiz





NO-SQL (CONTINUED)

NoSQL Taxonomy



- Key/Value
- Document
- Column
- Graph
- Others
 - Geospatial
 - File System
 - Object

Key-Value NoSQL Stores

Extremely simple interface

- Data model: (key, value) pairs
- Operations: Insert(key,value), Fetch(key), Update(key), Delete(key)
- Some allow (non-uniform) columns within value
- Some allow Fetch on range of keys

Implementation: efficiency, scalability, fault-tolerance

- Records distributed to nodes based on key
- Replication
- Single-record transactions, "eventual consistency"

Key-Value NoSQL: **Voldemort**

- ◉ Created by LinkedIn, now open source
- ◉ Inspired by Amazon's Dynamo
- ◉ Written in Java
- ◉ Pluggable Storage
 - BerkeleyDB, In Memory, MySQL
- ◉ Pluggable Serialization
 - JSON, Thrift, Protocol Buffers, etc.
- ◉ Cluster Rebalancing
- ◉ Versioning, based on Vector Clocks
 - Reconciliation occurs on reads.
- ◉ Partitioning and Replication based on Dynamo
 - Consistent Hashing
 - Virtual Nodes
 - Gossip

Document Stores

Like Key-Value Stores except value is document

- **Data model:** (key, document) **pairs**
- **Document:** JSON, XML, other semistructured formats
- **Basic operations:** Insert(key,document), Fetch(key), Update(key), Delete(key)
- Also Fetch based on document contents

Example systems

- CouchDB, MongoDB, SimpleDB, ...

CouchDB JSON Example



```
{
  "_id": "guid goes here",
  "_rev": "314159",

  "type": "abstract",

  "author": "Keith W. Hare"

  "title": "SQL Standard and NoSQL Databases",

  "body": "NoSQL databases (either no-SQL or Not Only SQL)
          are currently a hot topic in some parts of
          computing.",
  "creation_timestamp": "2011/05/10 13:30:00 +0004"
}
```


Document NoSQL: MongoDB

- Development started in 2007
- Commercially supported and developed by 10Gen
- Stores documents using BSON
- Supports AdHoc queries
 - Can query against embedded objects and arrays
- Support multiples types of indexing
- Officially supported drivers available for multiple languages
 - C, C++, Java, Javascript, Perl, PHP, Python and Ruby
- Community supported drivers include:
 - Scala, Node.js, Haskell, Erlang, Smalltalk
- Replication uses a master/slave model
- Scales horizontally via sharding
- Written C++



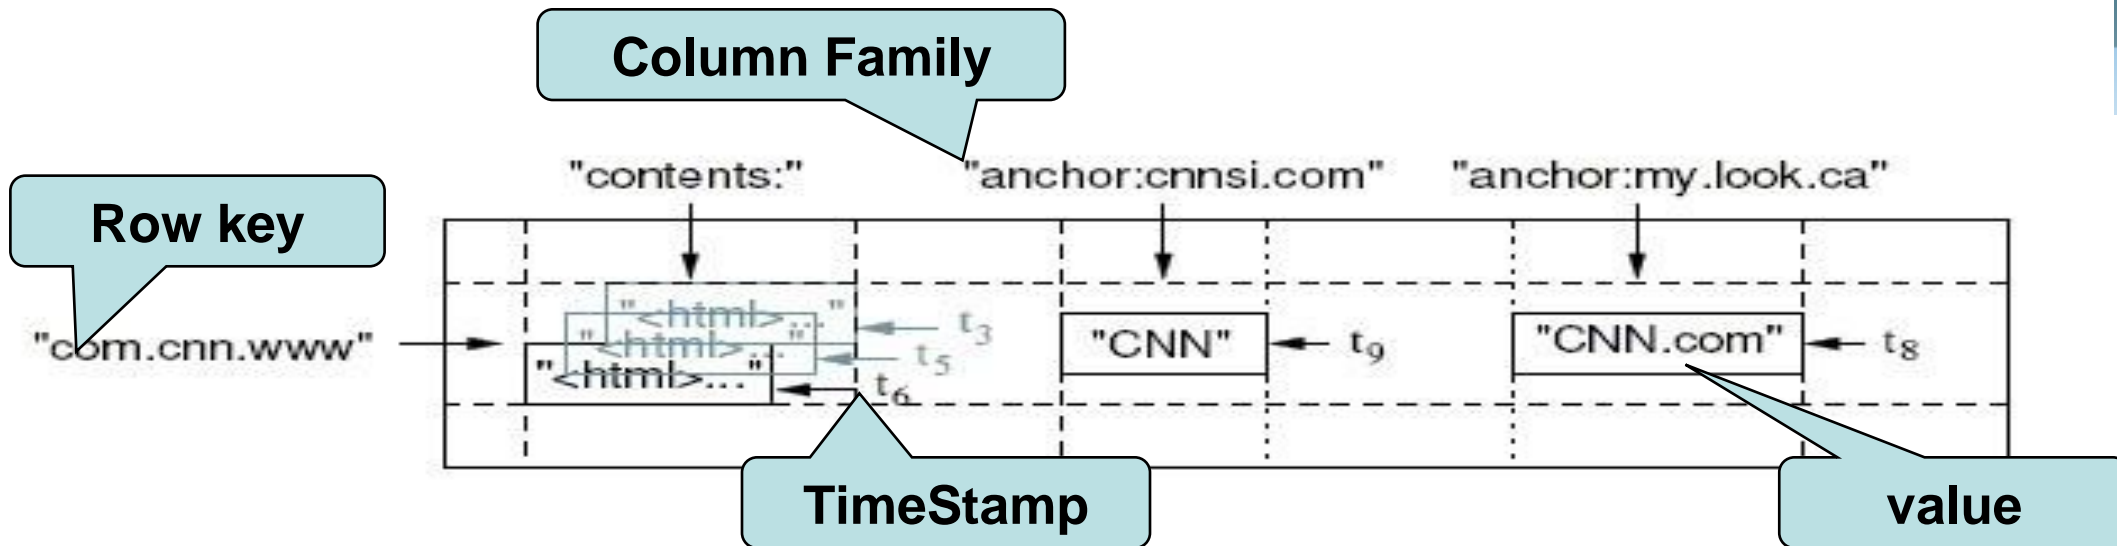
No-SQL

Google

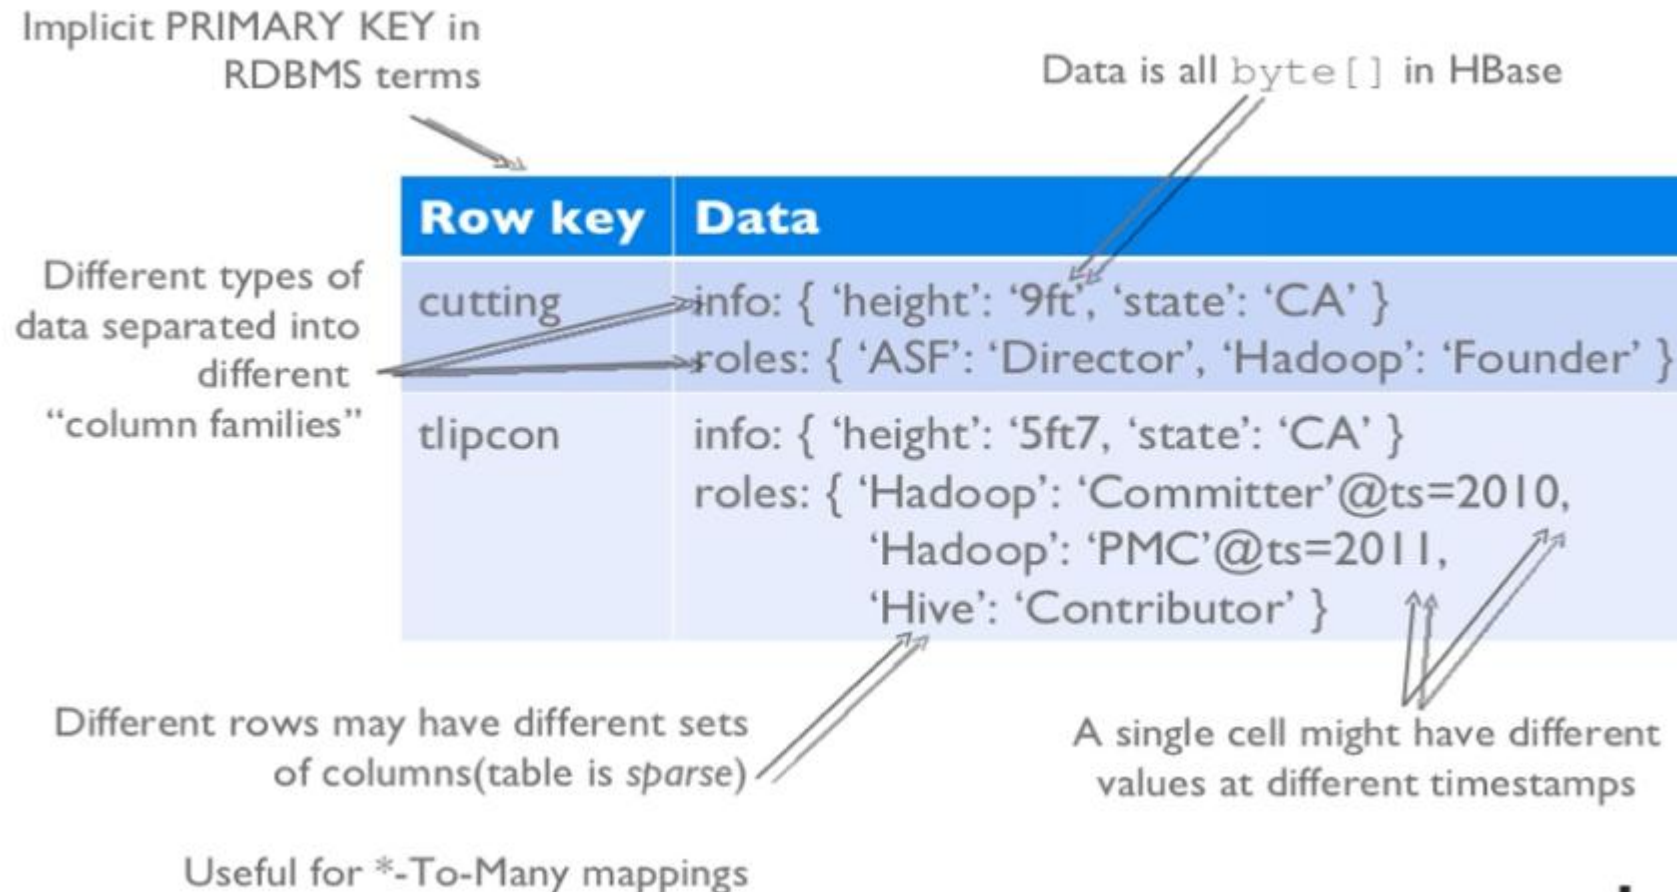
An **open-source**, **distributed**, **column-oriented** database built on top of HDFS based on **BigTable**!

Big Table: Data Model

- Tables are sorted by Row
- Table schema only defines its *column families*.
 - Each family consists of any number of columns
 - Each column consists of any number of versions
 - Columns only exist when inserted, NULLs are free.
 - Columns within a family are sorted and stored together
- Everything except table names are byte[]
- (Row, Family: Column, Timestamp) → Value



HBase Logical View



Physical distribution of the table

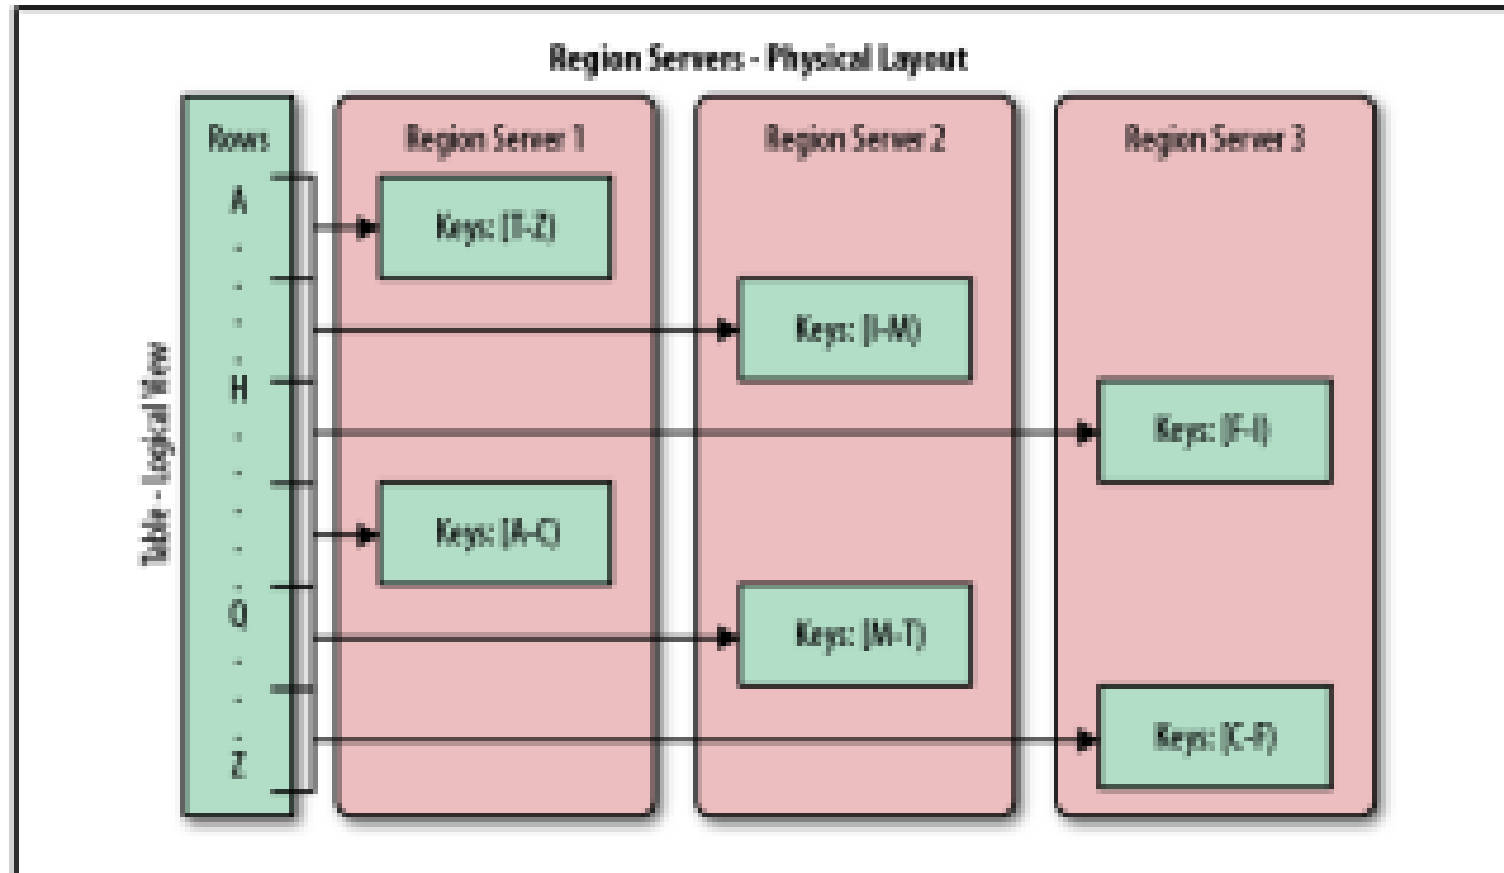
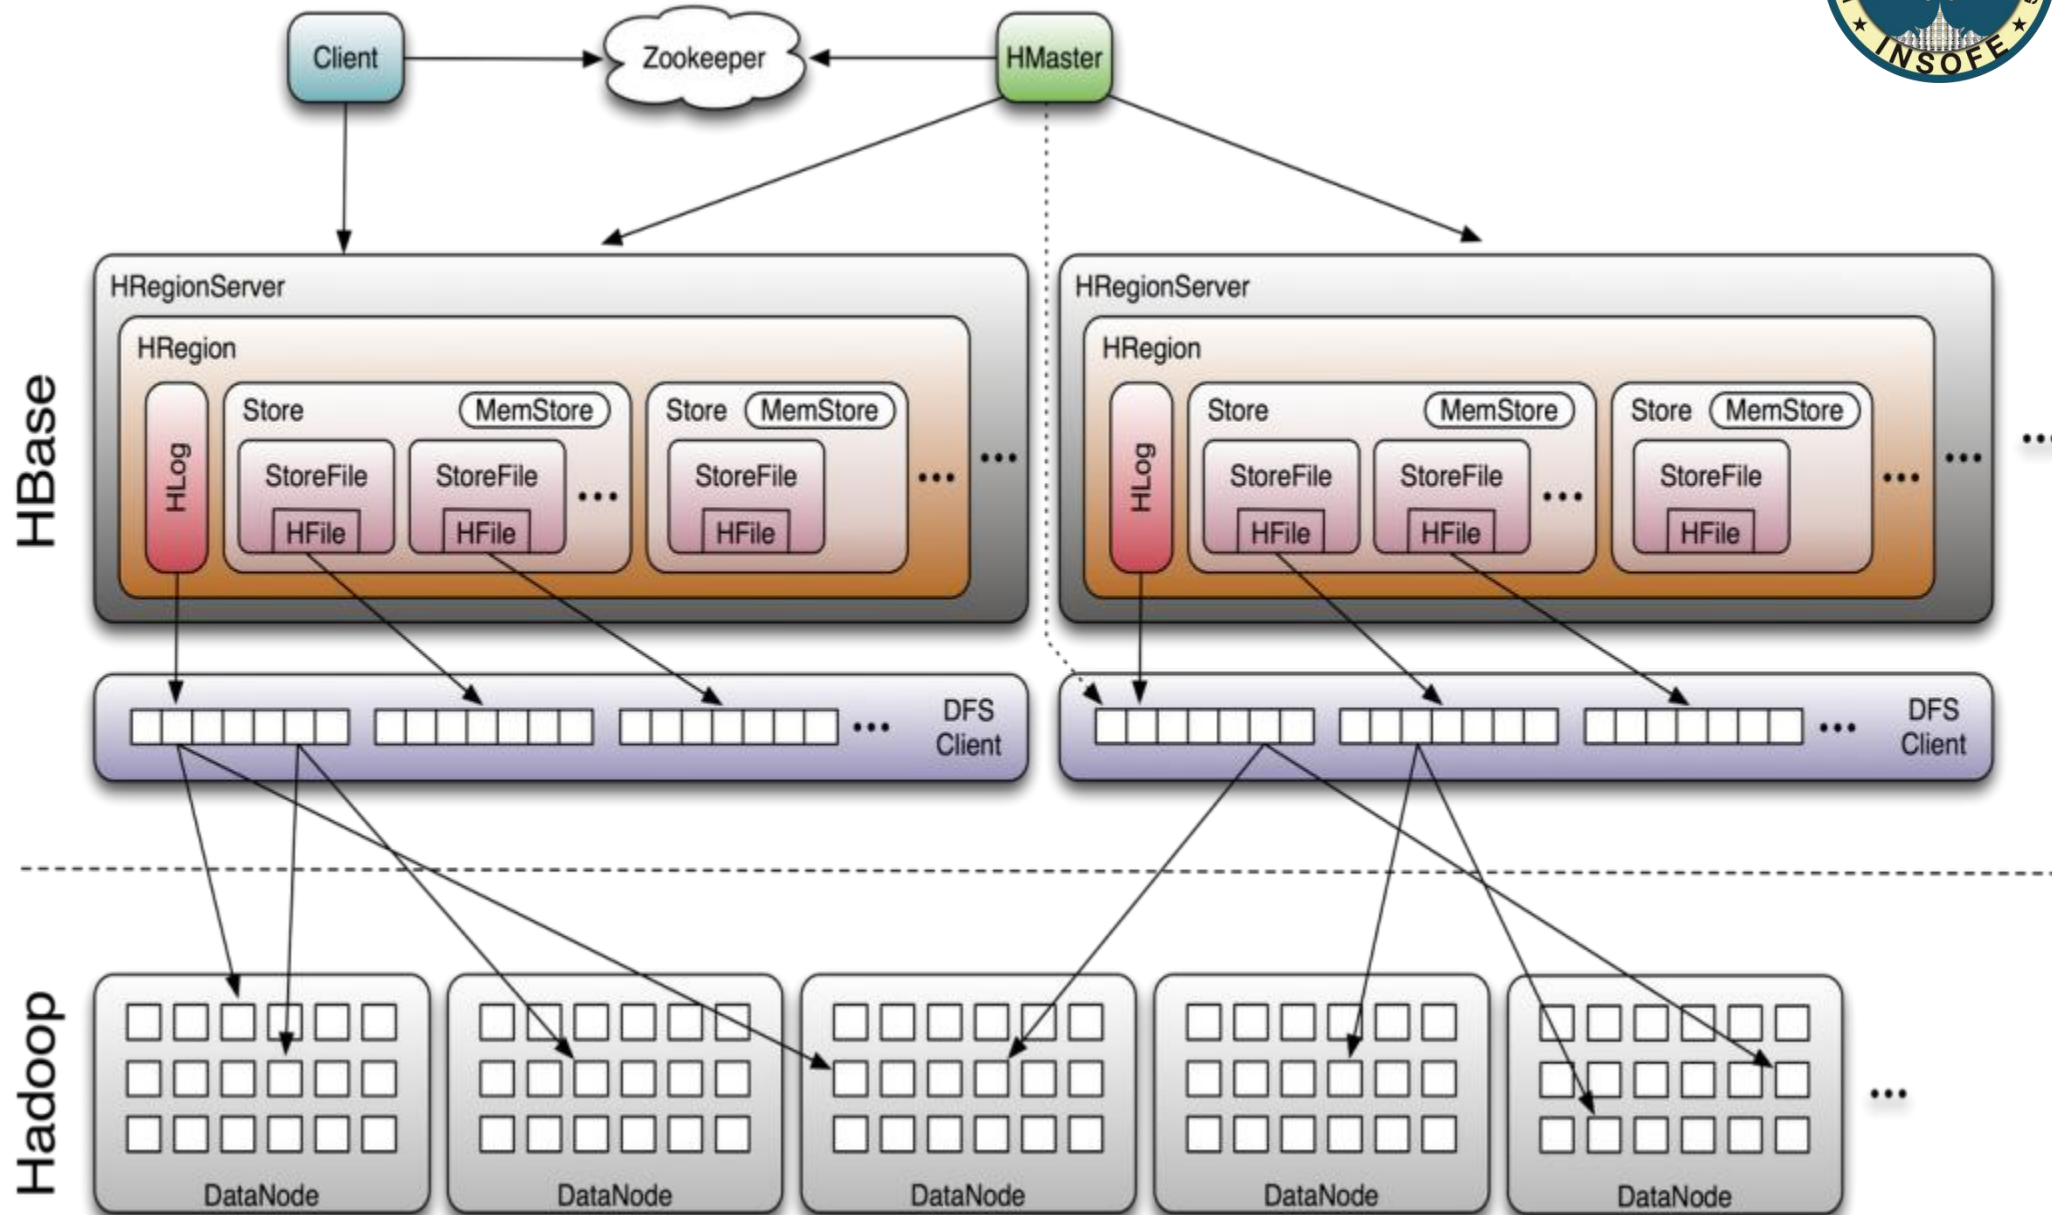


Figure 1-7. Rows grouped in regions and served by different servers

HBase implements BigTable on Hadoop



Write path of HBase

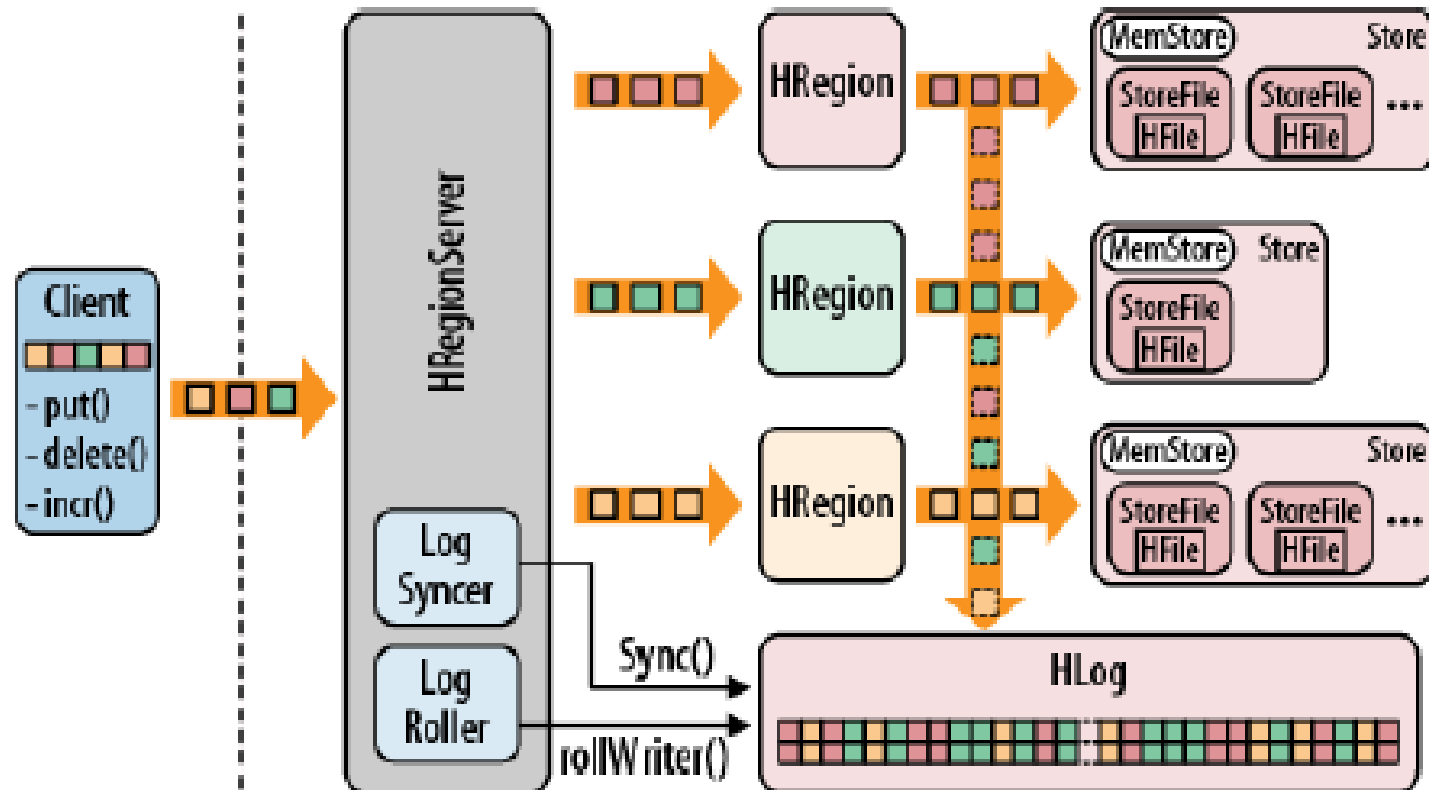


Figure: The write path of HBase

HBase vs. HDFS



	Plain HDFS/MR	HBase
Write pattern	Append-only	Random write, bulk incremental
Read pattern	Full table scan, partition table scan	Random read, small range scan, or table scan
Hive (SQL) performance	Very good	4-5x slower
Structured storage	Do-it-yourself / TSV / SequenceFile / Avro / ?	Sparse column-family data model
Max data size	30+ PB	~1PB

HBase vs. RDBMS



	RDBMS	HBase
Data layout	Row-oriented	Column-family-oriented
Transactions	Multi-row ACID	Single row only
Query language	SQL	get/put/scan/etc *
Security	Authentication/Authorization	Work in progress
Indexes	On arbitrary columns	Row-key only
Max data size	TBs	~1PB
Read/write throughput limits	1 000s queries/second	Millions of queries/second

When to use HBase

- You need random write, random read, or both (*but not neither*)
- You need to do many thousands of operations per second on multiple TB of data
- Your access patterns are well-known and simple



NoSQL: Performance Comparison

- Facebook Search
- MySQL > 50 GB Data
 - Writes Average : ~300 ms
 - Reads Average : ~350 ms
- Rewritten with Cassandra > 50 GB Data
 - Writes Average : 0.12 ms
 - Reads Average : 15 ms

Summary



- SQL Databases
 - Predefined Schema
 - Standard definition and interface language
 - Tight consistency
 - Well defined semantics
- NoSQL Database
 - No predefined Schema
 - Per-product definition and interface language
 - Getting an answer quickly is more important than getting a correct answer



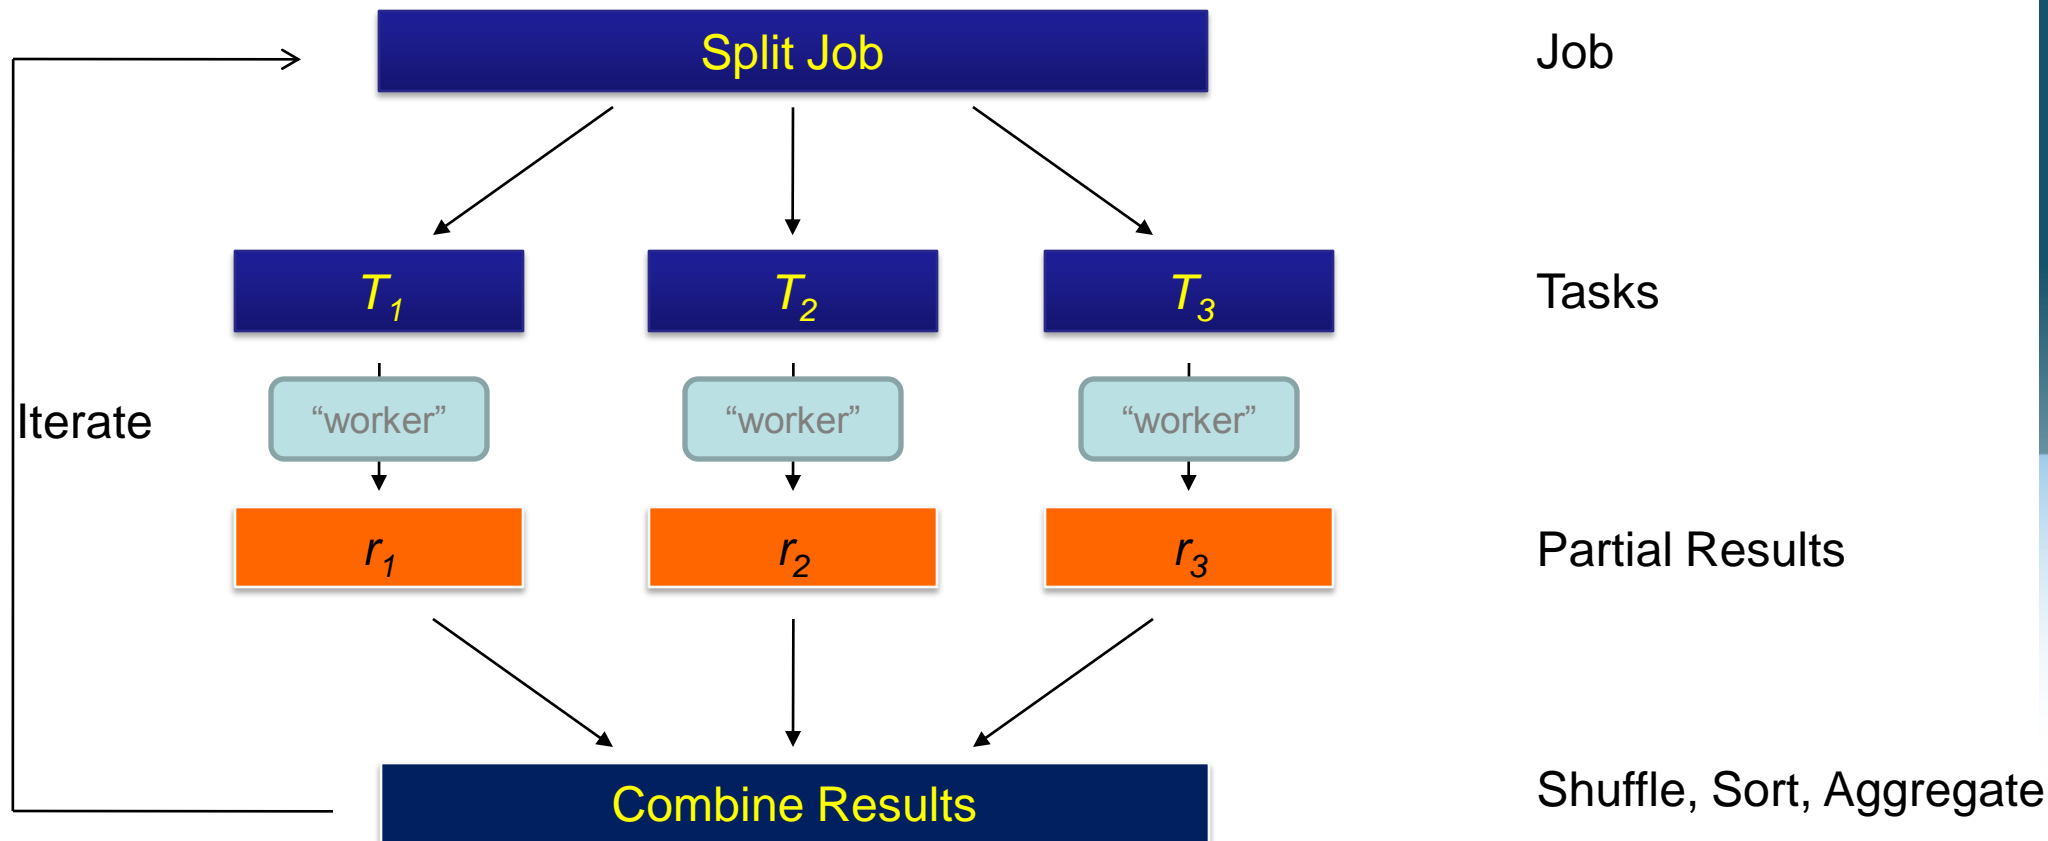
An excellent summary book on select NOSQL databases

WWW.CHRISTOF-STRAUCH.DE/NOSQLDBS.PDF



MAP REDUCE

Processing Frameworks 1: **Map-Reduce**



If not done right...



Source: Ricardo Guimarães Hermann



Map-Reduce: Solving large data problems


Map

- Iterate over a large number of records
- Extract something of interest from each
- Shuffle and sort intermediate results
- Aggregate intermediate results
- Generate final output

Reduce

Key idea: provide a functional abstraction for these two operations

Map-Reduce: Physical Flow



The Hello World of MapReduce

- Count the number of occurrences of each word in a large amount of input data
 - This is the 'hello world' of MapReduce programming

```
map(String input_key, String input_value)
    foreach word w in input_value:
        emit(w, 1)
```

```
reduce(String output_key,
        Iterator<int> intermediate_vals)
    set count = 0
    foreach v in intermediate_vals:
        count += v
    emit(output_key, count)
```

Hello World - continued

- Input to the Mapper:

```
(3414, 'the cat sat on the mat')  
(3437, 'the aardvark sat on the sofa')
```

- Output from the Mapper:

```
('the', 1), ('cat', 1), ('sat', 1), ('on', 1),  
( 'the', 1), ('mat', 1), ('the', 1), ('aardvark', 1),  
( 'sat', 1), ('on', 1), ('the', 1), ('sofa', 1)
```

- Intermediate data sent to the Reducer:


```
('aardvark', [1])  
( 'cat', [1])  
( 'mat', [1])  
( 'on', [1, 1])  
( 'sat', [1, 1])  
( 'sofa', [1])  
( 'the', [1, 1, 1, 1])
```

- Final Reducer output:

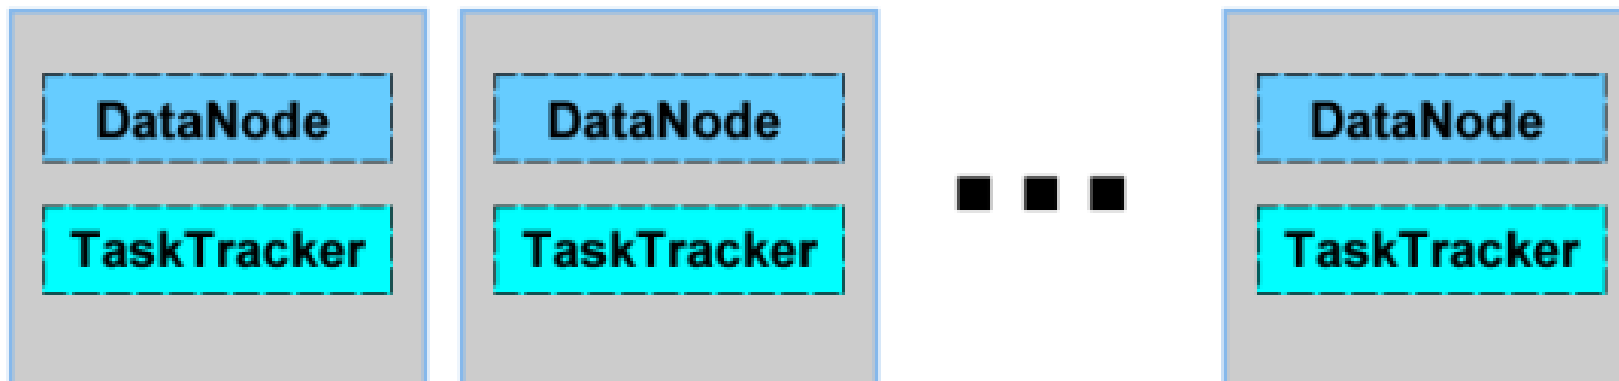
```
('aardvark', 1)  
( 'cat', 1)  
( 'mat', 1)  
( 'on', 2)  
( 'sat', 2)  
( 'sofa', 1)  
( 'the', 4)
```

Five Daemons of MapReduce

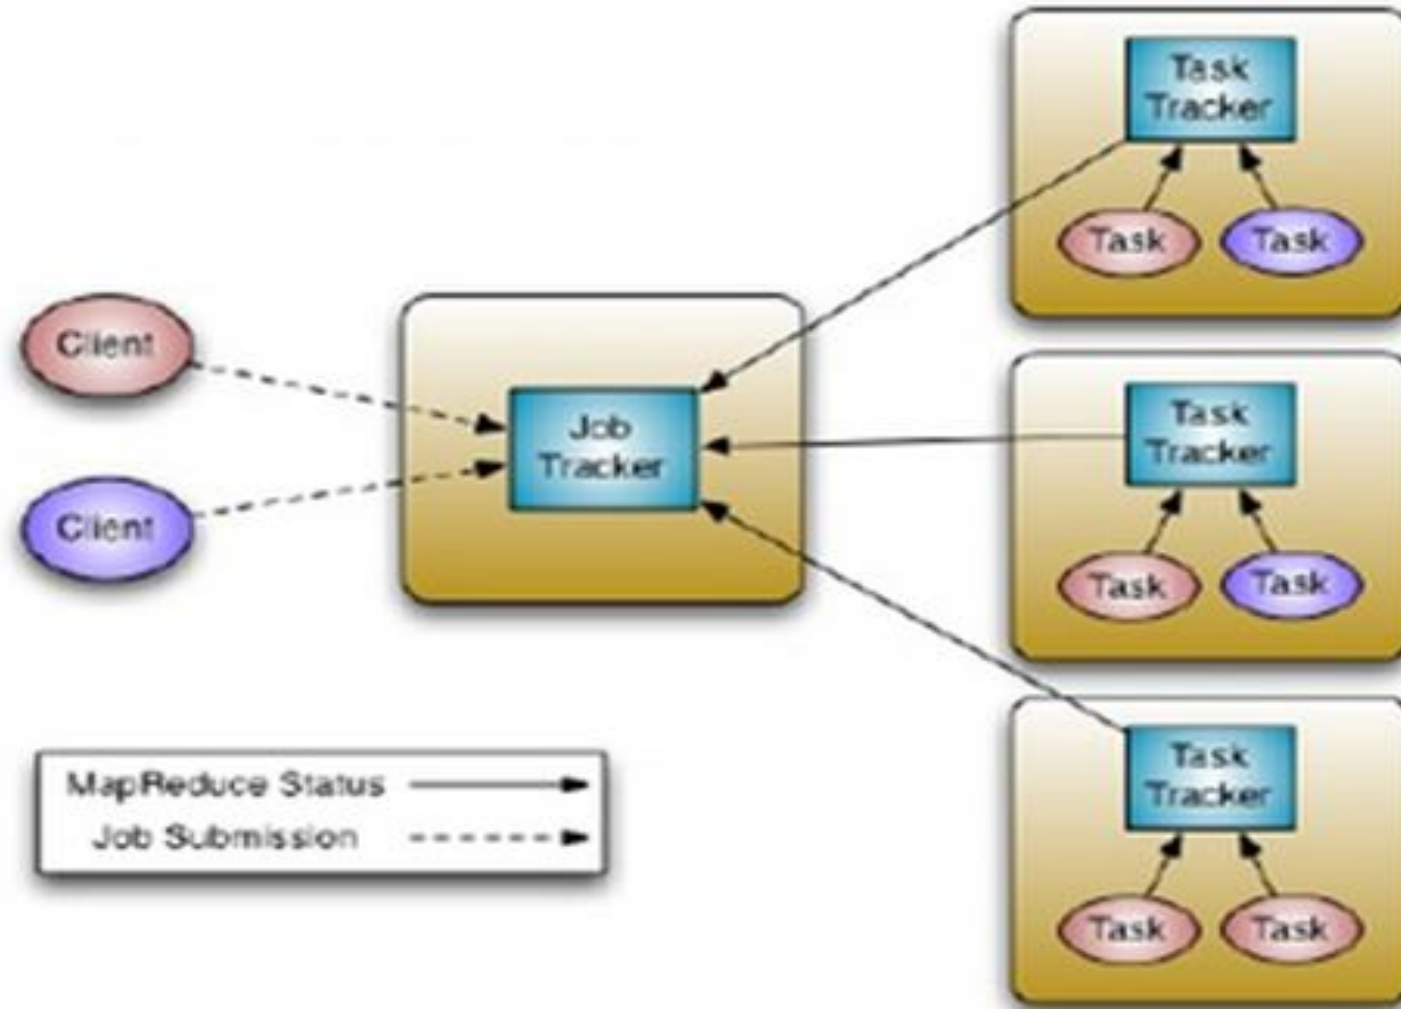
Master Nodes



Slave Nodes



Map Reduce Daemons (circa 2011)



Mapper

```
import java.io.IOException;

import org.apache.hadoop.io.IntWritable;
import org.apache.hadoop.io.LongWritable;
import org.apache.hadoop.io.Text;
import org.apache.hadoop.mapred.MapReduceBase;
import org.apache.hadoop.mapred.Mapper;
import org.apache.hadoop.mapred.OutputCollector;
import org.apache.hadoop.mapred.Reporter;

public class WordMapper extends MapReduceBase implements
    Mapper<LongWritable, Text, Text, IntWritable> {

    public void map(LongWritable key, Text value,
        OutputCollector<Text, IntWritable> output, Reporter reporter)
        throws IOException {
        String s = value.toString();
        for (String word : s.split("\\W+")) {
            if (word.length() > 0) {
                output.collect(new Text(word), new IntWritable(1));
            }
        }
    }
}
```

Reducer



```
import java.io.IOException;
import java.util.Iterator;

import org.apache.hadoop.io.IntWritable;
import org.apache.hadoop.io.Text;
import org.apache.hadoop.mapred.OutputCollector;
import org.apache.hadoop.mapred.MapReduceBase;
import org.apache.hadoop.mapred.Reducer;
import org.apache.hadoop.mapred.Reporter;

public class SumReducer extends MapReduceBase implements
    Reducer<Text, IntWritable, Text, IntWritable> {

    public void reduce(Text key, Iterator<IntWritable> values,
        OutputCollector<Text, IntWritable> output, Reporter reporter)
        throws IOException {

        int wordCount = 0;
        while (values.hasNext()) {
            IntWritable value = values.next();
            wordCount += value.get();
        }
        output.collect(key, new IntWritable(wordCount));
    }
}
```


The Driver Code

```
import org.apache.hadoop.fs.Path;
import org.apache.hadoop.io.IntWritable;
import org.apache.hadoop.io.Text;
import org.apache.hadoop.mapred.FileInputFormat;
import org.apache.hadoop.mapred.FileOutputFormat;
import org.apache.hadoop.mapred.JobClient;
import org.apache.hadoop.mapred.JobConf;
import org.apache.hadoop.conf.Configured;
import org.apache.hadoop.util.Tool;
import org.apache.hadoop.util.ToolRunner;

public class WordCount extends Configured implements Tool {
    public int run(String[] args) throws Exception {

        if (args.length != 2) {
            System.out.printf(
                "Usage: %s [generic options] <input dir> <output dir>\n",
                                                                    getClass().getSimpleName());
            ToolRunner.printGenericCommandUsage(System.out);
            return -1;
        }
        JobConf conf = new JobConf(getConf(), WordCount.class);
        conf.setJobName(this.getClass().getName());

        FileInputFormat.setInputPaths(conf, new Path(args[0]));
        FileOutputFormat.setOutputPath(conf, new Path(args[1]));

        conf.setMapperClass(WordMapper.class);
        conf.setReducerClass(SumReducer.class);
```

Driver code... contd

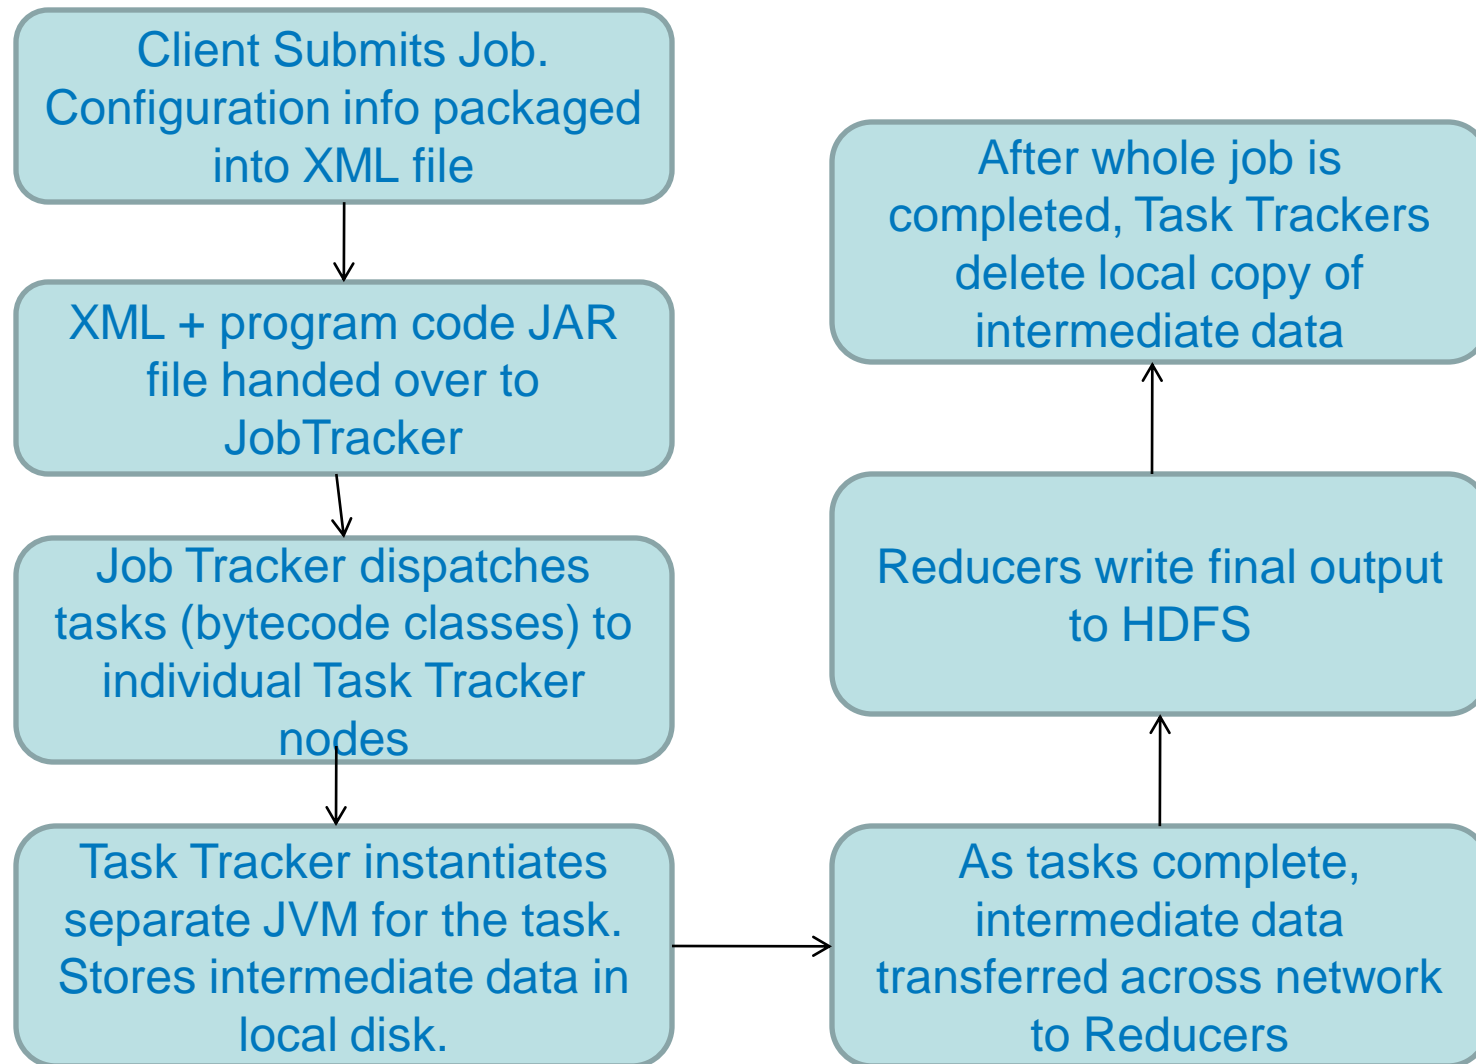
```
conf.setMapOutputKeyClass(Text.class);
conf.setMapOutputValueClass(IntWritable.class);

conf.setOutputKeyClass(Text.class);
conf.setOutputValueClass(IntWritable.class);

JobClient.runJob(conf);
return 0;
}

public static void main(String[] args) throws Exception {
    int exitCode = ToolRunner.run(new WordCount(), args);
    System.exit(exitCode);
}
}
```

Workflow



International School of Engineering

Plot 63/A, 1st Floor, Road # 13, Film Nagar, Jubilee Hills, Hyderabad - 500 033

For Individuals: +91-9502334561/63 or 040-65743991

For Corporates: +91-9618483483

Web: <http://www.insofe.edu.in>

Facebook: <https://www.facebook.com/insofe>

Twitter: <https://twitter.com/Insofeedu>

YouTube: <http://www.youtube.com/InsofeVideos>

SlideShare: <http://www.slideshare.net/INSOFE>

LinkedIn: <http://www.linkedin.com/company/international-school-of-engineering>